

INTRO TO THE HANDS-ON

- A public cloud platform: Google Cloud Platform
- A remote development environment : GitHub Codespaces

SDD - Data Engineering

<https://supaerodatascience.github.io/deep-learning/>



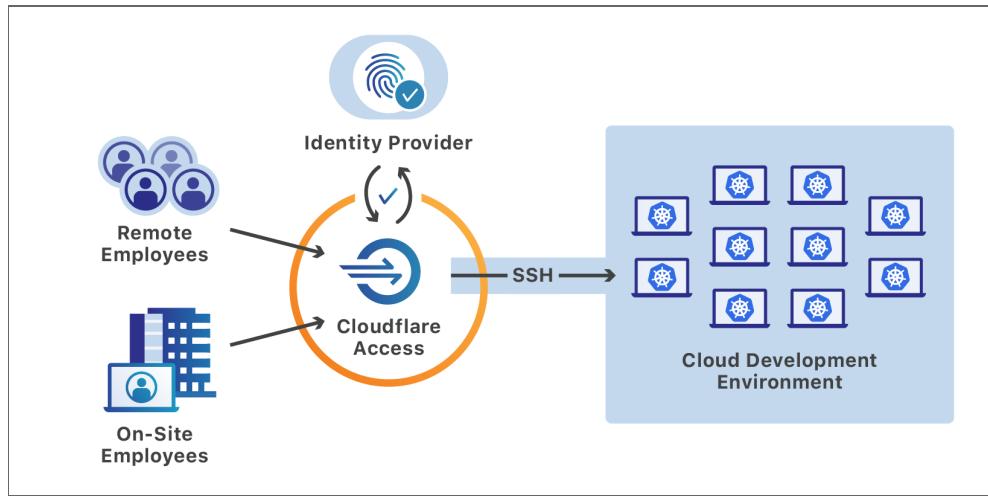
REMOTE DEVELOPMENT

WHY "REMOTE DEVELOPMENT" ?

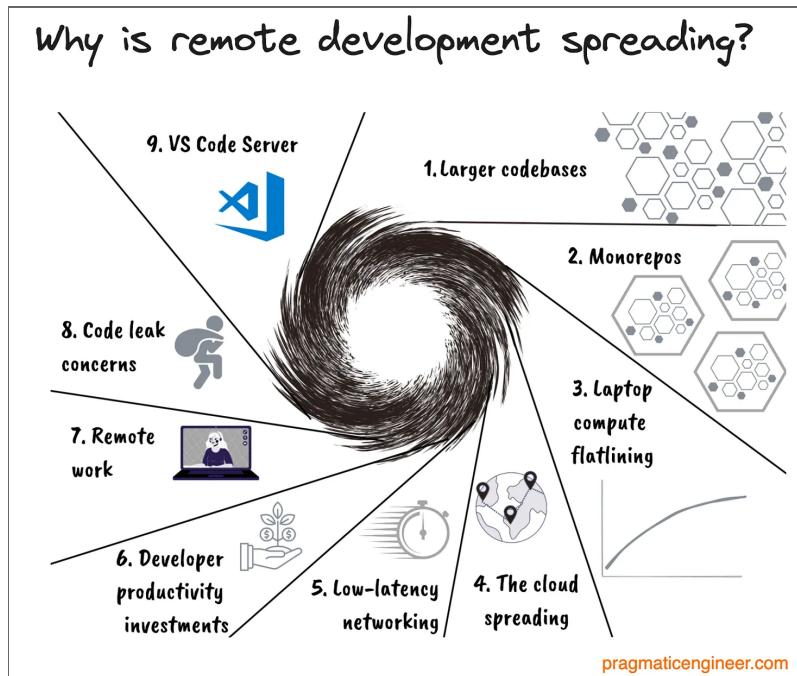
AI / Data Science = Data + Compute + Software

- Your job consist in handling huge volume of data
- Your job requires high computational resources
- You're working as a team with centralized computing platform
- You're working remotely

KEY USE CASE

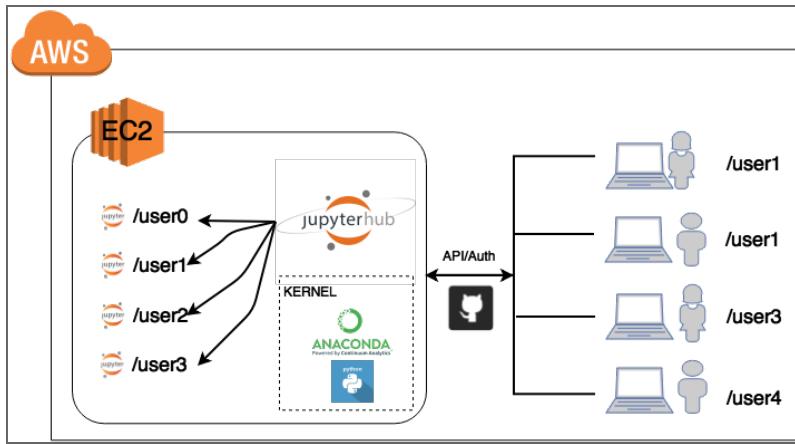


KEY USE CASE



<https://newsletter.pragmaticengineer.com/p/cloud-development-environments-why-now>

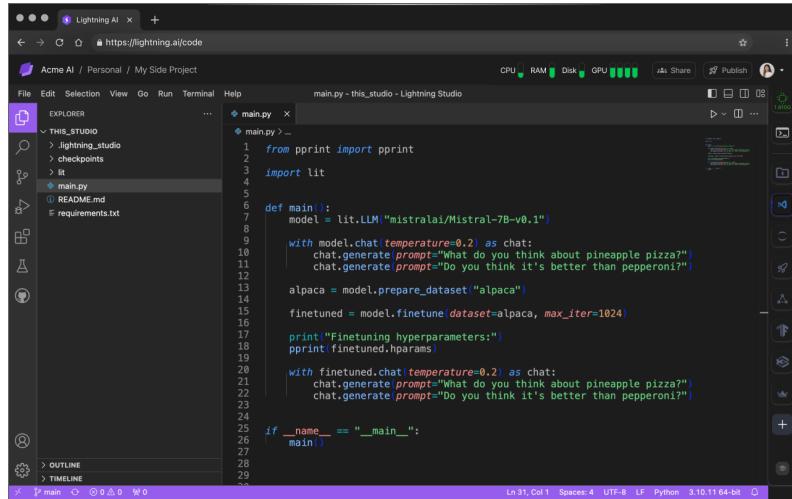
YOUR FUTURE DAILY ROUTINE



Another example (more managed) [Google Colab](#)

YOUR FUTURE DAILY ROUTINE

An AI-centric example, <https://lightning.ai/studios>



The screenshot shows the Lightning AI Studio interface. The main area is a code editor displaying a Python script named `main.py`. The code is used for finetuning a language model (mistralai/Mistral-7B-v0.1) on the Alpaca dataset. It includes imports for `pprint` and `lit`, defines a `main` function, and uses the `lit.LLM` class to initialize the model. The script then creates a `chat` object and generates responses to prompts about pizza toppings. It also prints finetuning hyperparameters and generates more responses. The code editor has syntax highlighting and line numbers. The top bar shows CPU, RAM, Disk, and GPU usage. The left sidebar shows project files like `README.md` and `requirements.txt`.

```
from pprint import pprint
import lit

def main():
    model = lit.LLM("mistralai/Mistral-7B-v0.1")

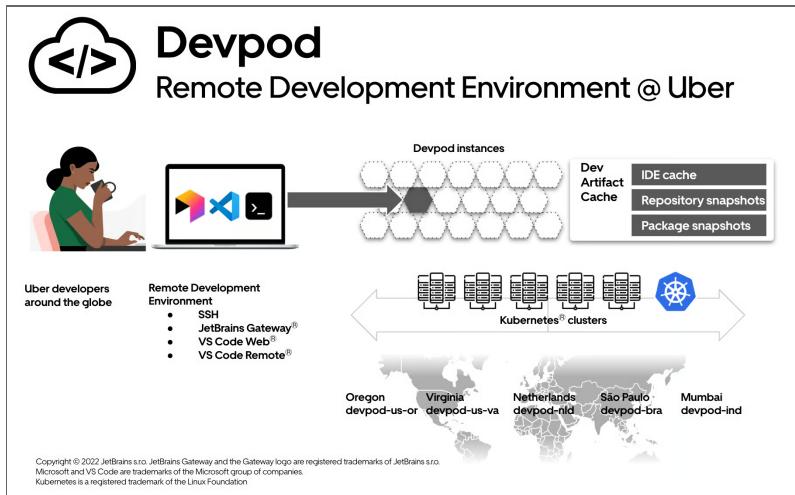
    with model.chat(temperature=0.2) as chat:
        chat.generate(prompt="What do you think about pineapple pizza?")
        chat.generate(prompt="Do you think it's better than pepperoni?")

    alpaca = model.prepare_dataset("alpaca")
    finetuned = model.finetune(dataset=alpaca, max_iter=1024)
    print("finetuning hyperparameters:")
    pprint(finetuned.hparams)

    with finetuned.chat(temperature=0.2) as chat:
        chat.generate(prompt="What do you think about pineapple pizza?")
        chat.generate(prompt="Do you think it's better than pepperoni?")

if __name__ == "__main__":
    main()
```

YOUR FUTURE DAILY ROUTINE



[Uber Blog describing their way of working](#)

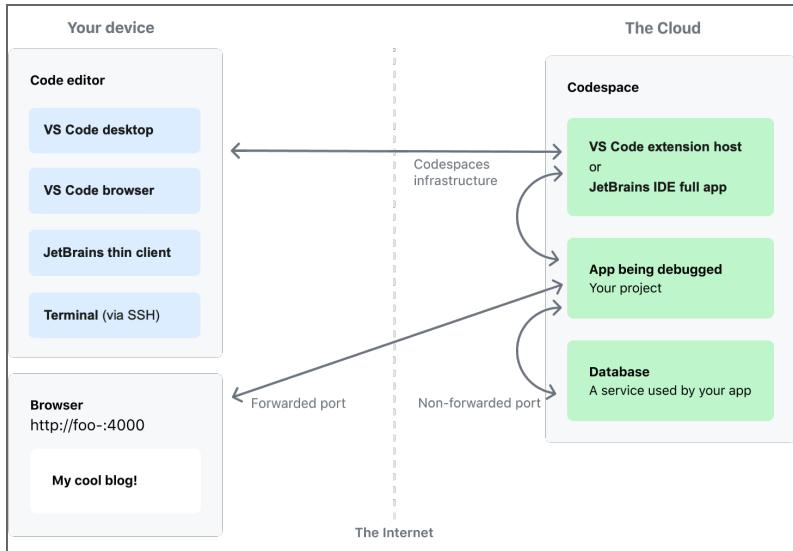
PROBLEMATICS

- How to transfer code ?
- How to interact with the machines ?
- How to get access to the data ?



GITHUB CODESPACES

- **Github Codespaces** : A managed development environment by Microsoft Azure
- A virtual machine and a **containerized development environment**
- A lot of built-in bonuses including "in-browser" connection & TCP port forwarding with reverse proxy

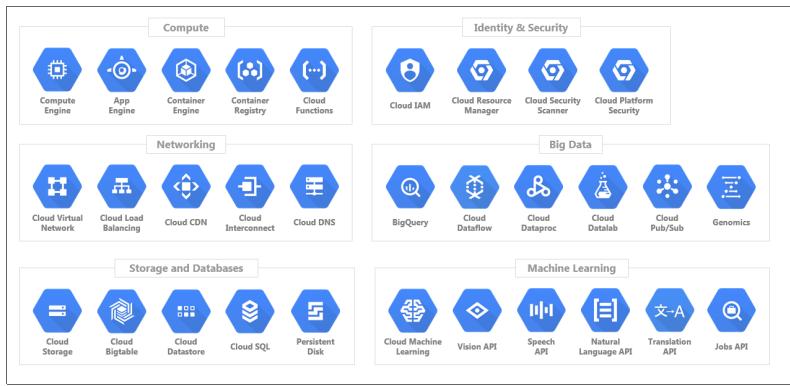


GOOGLE CLOUD PLATFORM



Google Cloud Platform

- One of the main cloud provider
- Behind AWS in SaaS (serverless...)
- More "readable" product line (for a Cloud Provider...)
- Very good "virtual machine" management
 - per second billing
 - fine-grained resource allocation

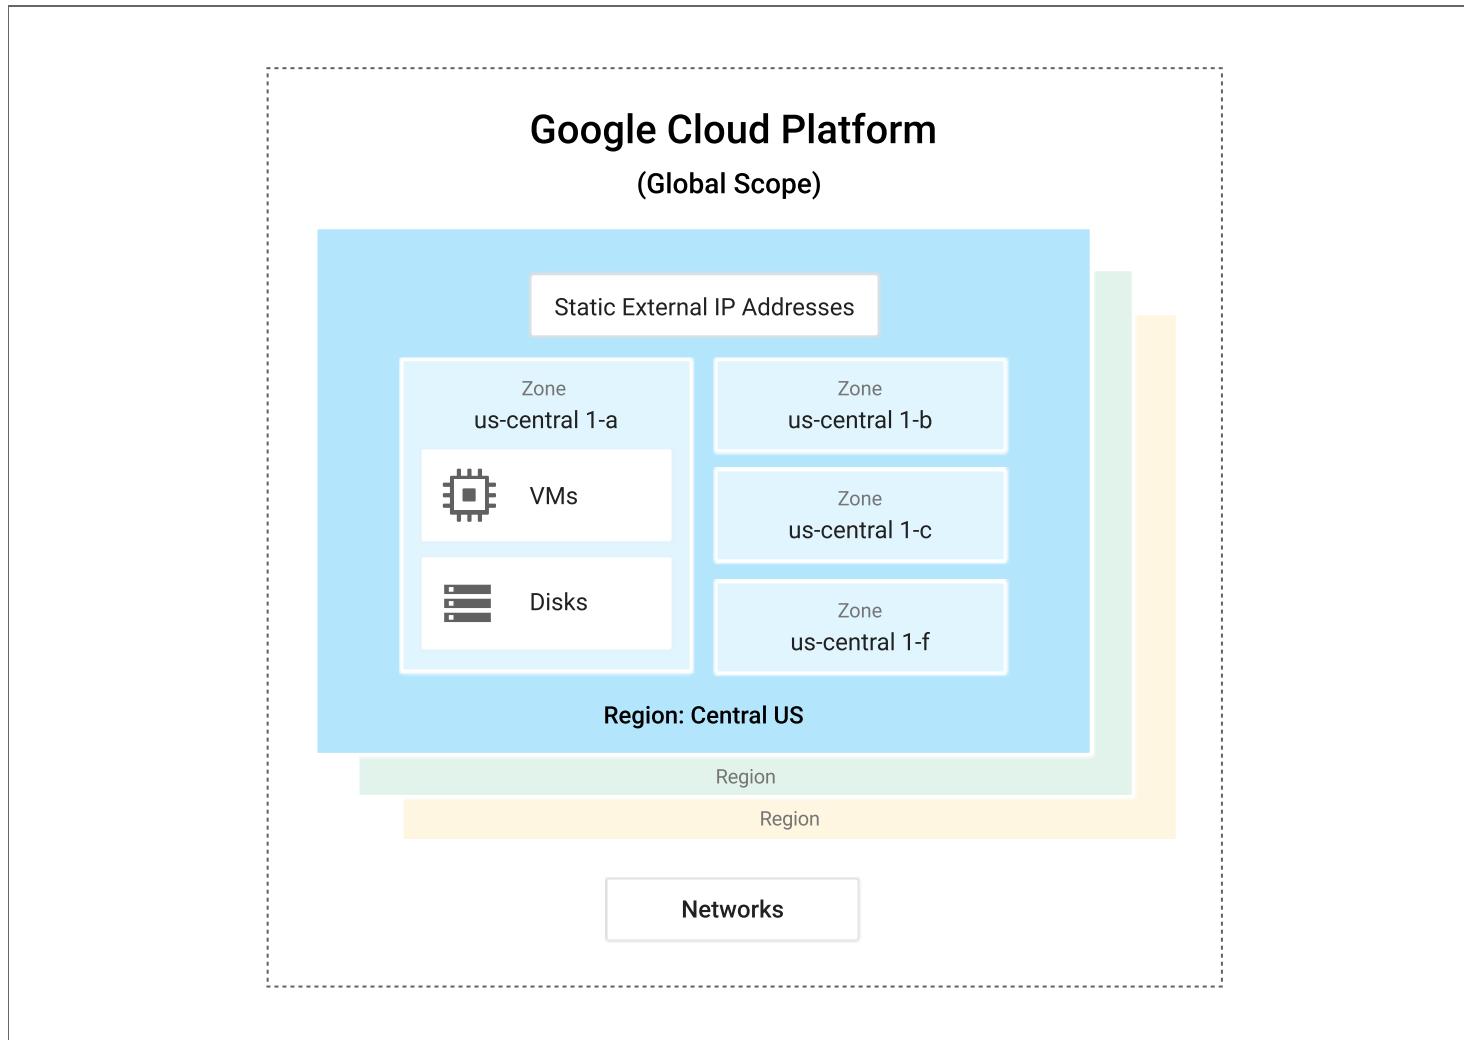




GCP Services in 4 words or less

CONCEPTS

ZONES AND REGIONS

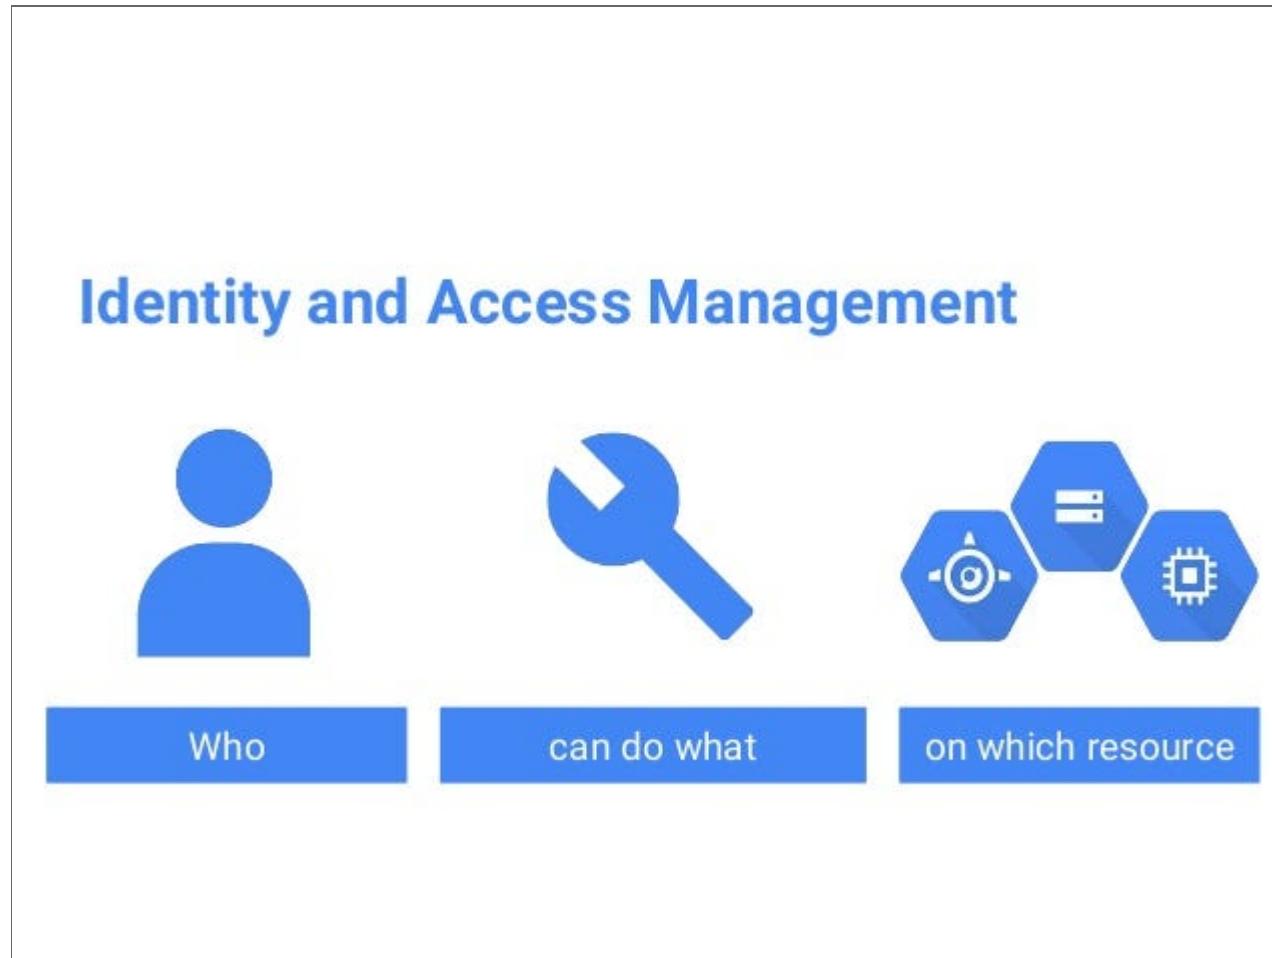


PROJECTS

The screenshot shows the 'PROJECTS' section of the Google Cloud Platform interface. At the top, it displays the project name 'Project: Example Project' and its ID 'ID: example-id (#123456789012)'. Below this, under the 'Resources' heading, there is a 'Storage' section. It shows a 'Cloud Storage' icon followed by the text '26 buckets'. At the bottom of the main content area, there is a link labeled 'Explore other services'.

- Access (Enabling API/Services)
- Resources (Quota by project)
- Networking
- Billing

CONCEPTS: IDENTITY AND ACCESS MANAGEMENT (IAM)



MAIN PRODUCTS WE ARE GOING TO BE LOOKING AT

- Google Compute Engine (virtual machine solutions)
- Google Cloud Storage (storage solutions)

GOOGLE COMPUTE ENGINE (GCE)

- The VM solution for GCP
- Images: Boot disks for VM instances example: **ubuntu-1804**
- Machine Types: Ressources available to your instance example: **n1-standard-8** (8 vCPU, 30 Gb RAM)
- Storage Options: "Attached disk" that can persist once the instance is destroyed... can be HDD, SSD...
- Preemptible: "Spot instances" on AWS", cheap but can be killed any minute by GCP

GOOGLE CLOUD STORAGE (GCS)

- Cheaper storage than persistent disks
- Can be shared between multiple instances / zones
- Higher latency than local disk
- Data is stored in "buckets" **whose name are globally unique**

GCS STORAGE CLASSES

Class	Use Case	Price
Standard	Frequent access	~\$0.02/GB
Nearline	Monthly access	~\$0.01/GB
Coldline	Quarterly access	~\$0.004/GB
Archive	Yearly access	~\$0.0012/GB
Choose based on access frequency!		

GCS WITH GSUTIL CLI

```
# List buckets  
gsutil ls  
  
# Create bucket (name must be globally  
unique!)  
gsutil mb gs://my-unique-bucket-name  
  
# Upload file  
gsutil cp local_file.csv gs://bucket/path/  
  
# Download file  
gsutil cp gs://bucket/path/file.csv ./  
  
# Recursive copy (for folders)  
gsutil -m cp -r gs://bucket/data/ ./  
local_data/
```

GCS ACCESS FROM VMS

VMs need **access scopes** to use GCS:

```
# When creating VM, add storage scope
gcloud compute instances create my-vm \
    --scopes=storage-rw # read-write access
to GCS
```

Or use a service account with appropriate IAM roles

INTERACTING WITH GCP: THE CONSOLE

The screenshot shows the Google Cloud Platform console dashboard. At the top left, it displays the project name 'OverviewExampleProject'. The main area is divided into several sections:

- Project Info:** Shows the project name, ID, and number, along with links to 'Go to project settings'.
- Resources:** Shows 1 Compute Engine instance and 1 BigQuery dataset.
- Trace:** Shows 'No trace data from the past 7 days'.
- Getting Started:** Includes links for enabling APIs, deploying a prebuilt solution, adding dynamic logging, monitoring errors with Error Reporting, deploying a Hello World app, creating a Cloud Storage bucket, creating a Cloud Function, and installing the Cloud SDK.
- Compute Engine:** Shows CPU usage over time (CPU (%)) with a chart peaking at approximately 0.8% around 2:45 PM.
- APIs:** Shows API requests per second (Requests (request/sec)) with a chart showing two peaks of approximately 0.14 requests/sec around 2:45 PM and 3:00 PM.
- Google Cloud Platform status:** Shows 'All services normal' with a link to 'Go to Cloud status dashboard'.
- Billing:** Shows estimated charges of USD \$0.21 for the billing period Jan 1 – 18, 2018, with a link to 'View detailed charges'.
- Error Reporting:** Shows 'No sign of any errors. Have you set up Error Reporting?' with a link to 'Learn how to set up Error Reporting'.
- News:** Lists recent news items:
 - 'Whitpaper: Embark on a journey from monoliths to microservices' (7 hours ago)
 - 'Analyzing your BigQuery usage with Ocado Technology's GCP' (5 days ago)
 - 'Running dedicated game servers in Kubernetes Engine: tutorial' (2 days ago)With a link to 'Read all news'.
- Documentation:** Links to 'Learn about Compute Engine', 'Learn about Cloud Storage', and 'Learn about App Engine'.

<https://console.cloud.google.com>

INTERACTING WITH GCP: SDK & CLOUD SHELL

- Using the gcloud CLI: <https://cloud.google.com/sdk/install>
- Using Google Cloud Shell: A small VM instance you can connect to with your browser
(that we won't use)

SELF-PACED HANDS-ON



OBJECTIVES

-  Discover GitHub Codespace
-  Create your GCP account, configure your credentials
- Create your first VMs and connect to it via SSH
 - Learn to use Tmux for detachable SSH sessions
 - Interact with Google Cloud Storage
 - End to End example :
 - Write code in your codespace
 - Transfer it to your GCE instance
 - Run the code, it produces a model
 - Transfer the model weights to google cloud storage
 - Get it back from your codespace
- Bonus content
 - Managed SQL
 - Infrastructure as code

SSH TUNNEL, PORT FORWARDING

*In computer networking, **a port is a communication endpoint**. At the software level, within an operating system, a port is a logical construct that identifies a specific process or a type of network service. **A port is identified for each transport protocol and address combination** by a 16-bit unsigned number, known as the port number. The most common transport protocols that use port numbers are the Transmission Control Protocol (TCP) and the User Datagram Protocol (UDP).*

[Wikipedia](#)

EXAMPLES OF PROTOCOLS & USUAL PORTS

Examples

- SSH on port 22
- HTTP on port 80
- HTTPS on port 443

http apps can serve content over specific ports

Example

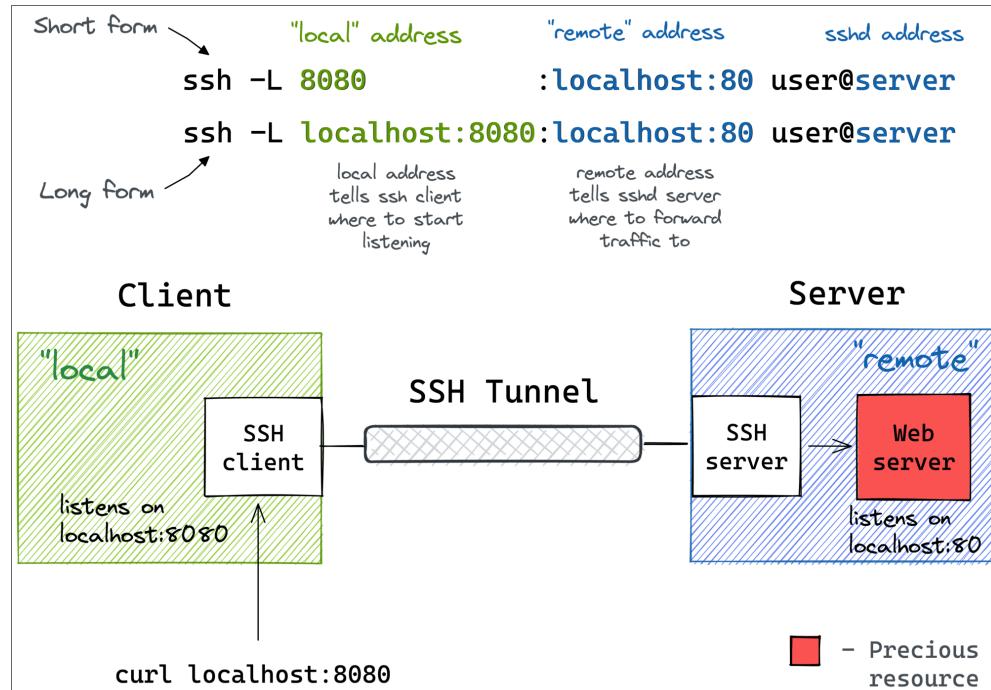
- Jupyter default is 8888 (that's why you open **http://localhost:8888**)

SSH TUNNELS

We usually connect to web app using `http://{ip}:{port}`

🤔 but what if the machine is not available from the public internet / local network ?

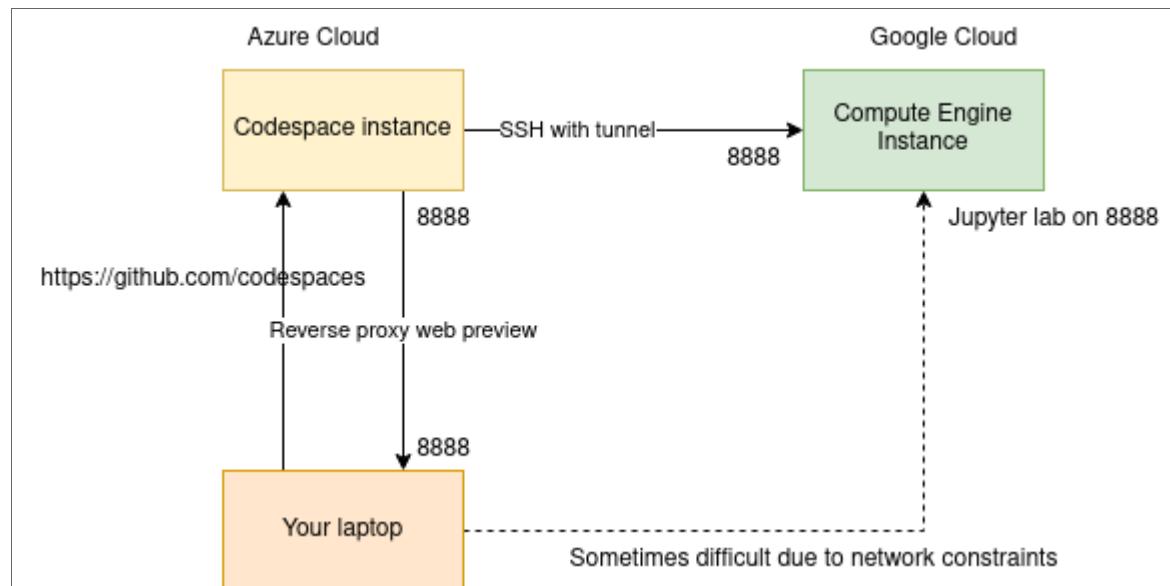
→ Enter SSH with port forwarding



Visual guide

TUNNELS OF TUNNELS

- Some of you did Local Machine -> (browser) -> Codespace -> (ssh) -> VM -> Jupyterlab on port 8888
- With port transfers !
- What happens when you go to **http://(url-generated-by-codespaces):8888** in this case ?



Speaker notes